

A convex-cone reachability oracle for cell-state engineering: falsifiable feasibility verdicts and constructive certificates from measured genome-scale perturbation screens

Cell-State Reachability Project¹

¹Built with Claude — Life Sciences (Research track)

July 10, 2026

Abstract

Deciding whether a desired change in cell state can be produced by a given class of perturbation—and, if not, saying so before the experiment is run—is the central question of cell-state engineering and of pharmaceutical target selection. Existing computational methods either predict the response to a specified perturbation or rank candidate regulators; none returns a *feasibility verdict* grounded in measured effects, and none can certify that a goal is unreachable. We introduce a convex-cone reachability oracle that treats a genome-scale CRISPRi Perturb-seq screen as a dictionary of measured effect vectors and asks whether a target cell-state shift lies in their non-negative combination. Because loss-of-function knockdowns combine additively and non-negatively, reachability is membership in a convex cone, decided by non-negative least squares; a reachable target yields a minimal ranked knockdown recipe, and an unreachable target yields a Farkas/KKT separating-hyperplane certificate that names the genes which would instead have to be *activated*—a falsifiable gain-of-function hypothesis rather than a null result. Applied to a screen in primary human CD4⁺ T cells, the oracle finds the T_H2→T_H1 polarization shift partly reachable by knockdown (held-out cosine 0.448 against a shuffled-target null of -0.005 , $z \approx 24$), with a signed decomposition attributing 39% of the shift to loss-of-function, 25% to required activation, and 35% to an irreducible remainder; canonical regulators appear as expected positive controls (GATA3 reached, TBX21 correctly anti-aligned). Across a 12-cell atlas of four transitions under three activation states, knockdown is never the majority modality (mean loss-of-function fraction 0.34). The same operator transfers unchanged to a K562 CRISPRa combinatorial screen (held-out CEBPA cosine 0.878), whose measured double perturbations show that the additivity assumption is a bounded approximation (median cosine 0.71 to the sum of singles) rather than an identity. Against a 92-method survey (2011–2026), measured-effect grounding and a feasibility verdict with a certificate are held by disjoint sets of prior methods; their intersection is empty. The verdict is stable under stress: it strengthens rather than weakens when the metric is reweighted by target magnitude to guard against signal dilution, the non-negativity constraint costs little in fit yet supplies the certificate, and a cross-cell-type probe shows that the *direction* of a shift transfers across cell types while the specific minimal *recipe* does not. The oracle is a hypothesis-prioritization and experiment-triage instrument—it decides where a screen should spend its arms and where it should stop—not a target-validation engine, and every nomination is a falsifiable wet-lab hypothesis.

1 Introduction

Much of experimental biology is a search for interventions that move a cell from one state to another: differentiating a stem cell, repolarizing a T cell, reverting a diseased phenotype, or—in drug discovery—knocking down a target to reverse a pathogenic program. Before committing reagents to such a program, the decisive question is not *what will this perturbation do* but *can the desired change be achieved at all with the perturbations available, and if not, what is missing*. That question is rarely asked computationally, and its economic weight is large: only about one in ten clinical programs succeeds, most failures are for lack of efficacy rather than safety, and the interventions with demonstrated power to shift those odds—human genetic support and directly measured target effects—both argue for grounding target choice in real perturbation data rather than inferred models [1–5].

The computational tools available to answer it fall into two camps that, as we document in Section 4, do not overlap. Perturbation-response models—from gene-regulatory-network simulators to single-cell foundation models—are increasingly trained on real screens but are forward predictors: given a perturbation they return a predicted response, and a predictor structurally always returns *some* answer, so it cannot certify that a goal is out of reach. Network-control and Boolean-attractor methods do reason about reachability and can carry controllability guarantees, but those guarantees are properties of a hand-built or inferred model, only as trustworthy as the network assumed. Neither camp answers the feasibility question from measured effects. This gap matters all the more because a wave of 2025–2026 benchmarks finds that current perturbation predictors, foundation models included, frequently fail to beat mean or linear baselines on genuinely unseen perturbations [6–9]: the field’s own evidence argues for an output weaker and more defensible than predicting an unseen response.

The convex-cone move. We take the opposite starting point from model building: we do not infer a network at all. A genome-scale CRISPRi Perturb-seq screen [10] directly measures, for each of tens of thousands of single-gene knockdowns, the effect vector it produces in expression space. We treat these measured vectors as a fixed dictionary and ask a geometric question: does the desired cell-state shift lie in the set of shifts the dictionary can produce? Two facts about knockdowns make this question exactly a convex-geometry problem. A knockdown can only remove function, fixing the sign of each atom’s contribution; and applying several knockdowns together produces, to first order, an additive combination with non-negative loadings. The set of achievable shifts is therefore the *non-negative conic hull* of the measured effect vectors, and asking whether a target is reachable is asking whether it lies inside that cone—a question non-negative least squares (NNLS) answers exactly [11]. This reframing turns an open-ended design search into a decision with a proof on both sides. When the target is inside the cone, the NNLS weights *are* a minimal ranked knockdown recipe. When it is outside, convex duality supplies a separating hyperplane—a Farkas/KKT certificate [12]—whose positive coordinates name the genes the target requires but that no combination of knockdowns can supply. An infeasibility result is thus not a dead end but a constructive, falsifiable hypothesis: it says the move requires *activation*, and it names the genes to activate, redirecting the experiment from a CRISPRi arm to a CRISPRa arm.

Scope and honesty. The oracle is a decision instrument, not an oracle of biological truth. It certifies feasibility *relative to the measured dictionary* under an explicit additivity model, and its verdicts are hypotheses to be tested at the bench, not validated targets. We are correspondingly conservative about what we claim. We report held-out-gene cosines against a shuffled-target null rather than in-sample fits, because the in-sample number is optimistic; we show that the absolute

reachable cosine of a real polarization shift is modest (0.448 for $T_H2 \rightarrow T_H1$), which is the honest ceiling of loss-of-function acting alone, not a defect of the method; and we verify on measured double perturbations that the additivity assumption underpinning multi-gene recipes is a bounded approximation rather than an identity. The regulators the oracle surfaces in T cells and in the CEBPA program are already reported by the source screens [10, 13]; we use them only as positive controls, and claim novelty for the measured-effect, certificate-carrying reachability *decision* itself.

Contributions. This paper makes four contributions.

1. **A feasibility oracle for cell-state engineering.** We formalize reachability of a target cell-state shift as membership in the convex cone of measured perturbation effect vectors, and solve it with NNLS so that a verdict, a minimal ranked recipe, and—on infeasibility—a Farkas/KKT activation certificate all fall out of one convex program (Section 2).
2. **A validated headline result with calibrated honesty.** On primary human $CD4^+$ T cells the oracle finds $T_H2 \rightarrow T_H1$ partly reachable by knockdown, far above a shuffled-target null but well below full reachability, with a signed loss-/gain-of-function decomposition and canonical regulators recovered as positive controls (Section 3.1).
3. **Generality across transitions and datasets.** A 12-cell atlas shows knockdown is never the majority modality across four transitions and three activation states, and the identical operator transfers without retuning to a K562 CRISPRa combinatorial screen, where we also quantify the additivity approximation directly on measured doubles (Sections 3.2–3.4).
4. **An exact positioning against the literature.** Over 92 methods spanning 2011–2026, we show that measured-effect grounding and feasibility-with-certificate are held by disjoint method sets whose intersection is empty, and translate the resulting verdicts into a GO/REDIRECT/STOP triage over 102 real target nominations (Sections 3.5–3.6).

2 Methods

2.1 Data and effect vectors

The primary dictionary is the genome-scale CRISPRi Perturb-seq screen of Zhu et al. [10] in primary human $CD4^+$ T cells (Marson and Pritchard laboratories, CZI Virtual Cells Platform), comprising differential-expression statistics for 33,983 single-gene knockdowns across 10,282 measured genes from four donors. For each knockdown j we take the published per-gene effect vector $\mathbf{E}_{:,j} \in \mathbb{R}^G$, the signed transcriptional shift the knockdown induces relative to non-targeting controls. Stacking these columns gives the effect matrix $\mathbf{E} \in \mathbb{R}^{G \times P}$. A *target* is a desired cell-state shift $\mathbf{d} \in \mathbb{R}^G$, computed as the difference between the mean expression profile of the source state and that of the destination state (e.g. T_H2 minus T_H1). All effect and target vectors are restricted to the genes measured in both the screen and the target contrast, and comparisons are reported as cosine similarity so that magnitude scaling does not affect the verdict. The identical pipeline is applied without modification to the K562 combinatorial CRISPR screen of Norman et al. [13] (Section 3.4); no dataset-specific parameters are tuned.

2.2 Reachability as membership in a convex cone

Two properties of loss-of-function perturbations determine the geometry. First, a CRISPRi knockdown removes rather than adds function, so each atom $\mathbf{E}_{:,j}$ enters a combination with a fixed

orientation. Second, applying a set of knockdowns together produces, to first order, the sum of their individual effects; there is no natural notion of a “negative amount” of a knockdown. The set of cell-state shifts reachable by a non-negative combination of knockdowns is therefore the *non-negative conic hull* of the measured effect vectors,

$$\mathcal{C} = \left\{ \mathbf{E}\mathbf{w} : \mathbf{w} \in \mathbb{R}_{\geq 0}^P \right\} \subseteq \mathbb{R}^G, \quad (1)$$

a closed convex cone. Deciding whether a target $\mathbf{d}$ is reachable is deciding whether $\mathbf{d} \in \mathcal{C}$ (up to a non-negative scale, since only direction is claimed). We quantify proximity by projecting $\mathbf{d}$ onto $\mathcal{C}$: the reachable component is the solution of the non-negative least-squares (NNLS) program

$$\mathbf{w}^* = \arg \min_{\mathbf{w} \geq 0} \|\mathbf{E}\mathbf{w} - \mathbf{d}\|_2^2, \quad \hat{\mathbf{d}} = \mathbf{E}\mathbf{w}^*, \quad (2)$$

solved with the Lawson–Hanson active-set algorithm [11]. The *reachable cosine* $\cos(\hat{\mathbf{d}}, \mathbf{d})$ measures how much of the target’s direction the cone can express, and the residual $\mathbf{d} - \hat{\mathbf{d}}$ is the part it cannot. A target is declared reachable when its reachable cosine is both high in absolute terms and far above the null of Section 2.5; the same solve returns the support of $\mathbf{w}^*$, whose largest entries constitute a *minimal ranked knockdown recipe*. Ordering genes by a greedy forward selection that adds at each step the knockdown giving the largest marginal gain in reachable cosine yields the reachability spectrum reported throughout.

2.3 Infeasibility certificates via convex duality

The value of a feasibility oracle is that “no” is provable, not merely observed. When $\mathbf{d} \notin \mathcal{C}$, the separating-hyperplane theorem for convex cones (a form of Farkas’ lemma [12]) guarantees a vector $\mathbf{y}$ with $\mathbf{E}^\top \mathbf{y} \leq 0$ and $\mathbf{d}^\top \mathbf{y} > 0$: a hyperplane with the whole cone on one side and the target strictly on the other. This $\mathbf{y}$ is exactly the negative gradient of the NNLS objective at the optimum, so it is produced by the same solve rather than a separate computation. The Karush–Kuhn–Tucker (KKT) stationarity conditions of (2) provide the numerical proof: at the optimum, active (positive-weight) atoms have zero reduced gradient and inactive atoms have non-positive reduced gradient. We report the maximum stationarity violation as the certificate residual; across all reported solves it is at machine precision ($\approx 1.1 \times 10^{-11}$), confirming true optimality rather than early stopping. The biological reading of the certificate is constructive: the positive coordinates of the residual direction—the part of the target orthogonal to everything the cone can supply—name genes the target requires to be *raised* that no knockdown combination can deliver. These become falsifiable CRISPRa (activation) hypotheses, turning an infeasibility result into a redirected experiment rather than a null one.

2.4 Signed loss-/gain-of-function decomposition

To separate “cannot be reached because it needs activation” from “cannot be reached for other reasons,” we decompose the target into three orthogonal budgets as fractions of its squared norm: the component captured by the non-negative knockdown cone (loss-of-function, LOF); the component lying along directions that require increased expression (gain-of-function, GOF), obtained by projecting the residual onto the positive-expression orthant; and an irreducible remainder aligned with neither (“neither”). We report this in two conventions. The *signed* decomposition partitions the target by the sign of each coordinate’s required change; the *one-sided* decomposition instead attributes the residual after the non-negative projection. For $T_{H2 \rightarrow H1}$ these give LOF/GOF/neither of

39%/25%/35% (signed) and 39%/31%/30% (one-sided). We report both because they answer subtly different questions and neither is privileged; the qualitative conclusion—that loss-of-function alone accounts for well under half of the shift—is common to both.

2.5 Held-out validation and null models

Two distinct null models guard against two distinct failure modes, and we are explicit about which each addresses.

Held-out-gene generalization. An in-sample reachable cosine can be inflated by a solver freely reweighting thousands of atoms to fit a fixed target; it measures expressivity of the dictionary, not predictive validity. We therefore report a *held-out-gene* cosine: we fit the NNLS weights using only a random subset of measured genes (rows of $\mathbf{E}$ and $\mathbf{d}$) and evaluate the reachable cosine on the held-out genes, so the recipe must generalize to expression coordinates it was not fit on. For $T_{H2} \rightarrow T_{H1}$ the held-out cosine is 0.448, against an in-sample value of 0.627; the gap is the honest cost of generalization, and we quote the held-out number as the headline throughout.

Shuffled-target null. To calibrate what reachable cosine arises by chance from a high-dimensional cone, we recompute the held-out reachability against 60 random targets generated by shuffling the gene labels of $\mathbf{d}$, preserving its marginal magnitude distribution while destroying its biological structure. For $T_{H2} \rightarrow T_{H1}$ the null has mean -0.005 , standard deviation 0.019, and maximum 0.029; the observed held-out cosine sits at $z \approx 24$ above it. This null answers “is the target special,” distinct from held-out-gene generalization, which answers “does the recipe transfer to unseen coordinates.”

Analytic anisotropy null. Because a shuffled-target null over an atlas of targets is expensive, we derive a closed-form approximation to the expected null cosine as a function of two geometric quantities—the alignment of the target with the dominant direction of the cone and the effective anisotropy of the atom cloud (Supplementary Fig. S1 legend). Validated against the shuffle null on both synthetic and real $CD4^+$ data, it reproduces the empirical null to Pearson $r = 0.995$ (in-sample) and 0.998 (held-out), with maximum absolute error 0.047 and 0.07 respectively, replacing on the order of a thousand target shuffles with ten to twenty NNLS solves. This makes the atlas-wide reliability analysis of Section 3.2 tractable while keeping the empirical shuffle null as the reference for the headline result.

Positive controls. We verify that the geometry recovers known biology using canonical T-helper regulators as controls, *not* as discoveries. On the toward- T_{H1} track, GATA3—the master T_{H2} transcription factor, which must be removed—is reached at rank 155/6871 (97.7th percentile, cosine +0.101) on the loss-of-function side, while TBX21—the master T_{H1} factor, which must be raised—is correctly anti-aligned under knockdown at rank 6775/6871 (1.4th percentile). Because these regulators are already reported by the source screen [10], their recovery validates the operator without being claimed as a novel finding.

2.6 Robustness and reinforcement analyses

Five analyses probe whether the headline verdict is an artifact of a specific modelling or evaluation choice. All reuse `reachability.py` without modification.

Non-negativity ablation. To test whether the biological non-negativity constraint earns its place, we compare the held-out cosine of the constrained NNLS solve against an unconstrained least-squares solve (which permits negative weights, i.e. treating a knockdown as if it could be run in reverse) and against the single best-aligned generator, across all 12 atlas cells. We report the cosine cost of the constraint, the number of biologically unrealizable negative weights the unconstrained solution requires, and the fraction of the unconstrained cosine the cone recovers.

Reachable-cosine ceiling. Because a knockdown-only cone cannot express gain-of-function directions, the achievable held-out cosine has a modality-imposed ceiling. We compute this ceiling as $\sqrt{f_{\text{LOF}}}$, where f_{LOF} is the loss-of-function fraction of the signed decomposition, and express each headline cosine as a fraction of its own ceiling so that a “modest” absolute number is reported against what is attainable in principle rather than against 1.

DEG-magnitude weighting. Uniform-gene metrics can dilute a real signal by averaging over thousands of unchanged genes, a critique raised for perturbation-response evaluation by Mejia et al. [14]. We recompute every reachable cosine and held-out z under a weighting $w_g = |d_g|$ that concentrates the metric on the genes the target actually moves, and we recalibrate the dynamic-range fraction against a shuffled-gene floor and an interpolated-duplicate ceiling under the same weighting. The unweighted metric remains the default and reproduces the published numbers exactly; the weighting is a stress test, not a re-definition.

Cross-cell-type transfer. To ask whether the operator generalizes beyond a single cell type, we apply it to the genome-scale essential-gene Perturb-seq screens of Replogle et al. [15] in K562 and RPE1 cells (843 shared single-gene perturbations over 2,832 shared genes). We evaluate transfer at three levels: whether a single perturbation’s effect *direction* agrees across cell types (matched cosine against a shuffled-gene null); whether the reachability *verdict* agrees (graded by the residual norm within- versus cross-basis, since a binary verdict is uninformative when an 843-atom cone is highly expressive); and whether the minimal *recipe* agrees (Jaccard overlap of selected atoms against a random-selection null).

Certificate stability and external cross-reference. We test the activation certificate two ways. For reproducibility, we refit the certificate on a random half of the gene axis and measure the rank agreement of the top genes against the full-data ranking and across independent halves (Spearman ρ over 12 random splits, against a label-shuffle null). For biological interpretation, we retrospectively cross-reference the top certificate genes against the literature direction-of-effect and against Open Targets genetic-association scores for the relevant T_H1-driven autoimmune diseases. This cross-reference is an exploratory annotation, not a pre-registered test, and we report it as such.

2.7 Software and reproducibility

The oracle is implemented in a single module (`reachability.py`) built on standard NNLS routines; all reported quantities derive from the published effect matrices of Zhu et al. [10] and Norman et al. [13] with no additional model fitting. Every numeric value in this manuscript is drawn from a locked facts sheet tied to the analysis outputs, and each figure is generated directly from those outputs; figures were checked for exact correspondence to the facts sheet before inclusion.

3 Results

Figure 1 summarizes the framework: a genome-scale screen yields measured effect vectors (1A); reachability is membership in their convex cone, with the residual naming what knockdown cannot supply (1B); a reachable target returns a verdict, a signed decomposition, and an activation certificate (1C); and the same operator transfers to an independent screen (1D).




Figure 1: The convex-cone reachability oracle. (A) A genome-scale CRISPRi Perturb-seq screen in primary human $CD4^+$ T cells yields, for each of 33,983 knockdowns, a measured effect vector over 10,282 genes. (B) Reachability is membership in the non-negative conic hull of these vectors; the reachable component is the NNLS projection of the target, and the residual is the activation-required direction, with “outside” certified by a separating hyperplane (Farkas/KKT). (C) For $T_H2 \rightarrow T_H1$ the oracle returns a *partly reachable* verdict (held-out cosine 0.448, $\|KKT/Farkas\| \approx 1.1 \times 10^{-11}$), a signed decomposition (39% LOF / 25% GOF / 35% neither), and an activation certificate naming genes that must be raised. (D) The identical operator, with no retuning, transfers to the Norman K562 CRISPRa screen (held-out CEBPA cosine 0.878).

3.1 The $T_H2 \rightarrow T_H1$ shift is partly reachable by knockdown

Our headline target is the $T_H2 \rightarrow T_H1$ polarization shift in resting cells, a canonical and therapeutically relevant repolarization. The greedy reachability spectrum (Fig. 2a) rises rapidly and plateaus: the reachable cosine reaches a knee at $k = 7$ knockdowns (cosine 0.40) and saturates near 0.45, sitting far above the shuffled-target null at every recipe size. The held-out-gene cosine is 0.448 against a null of mean -0.005 and standard deviation 0.019 ($z \approx 24$), so the target is unambiguously reachable in the statistical sense—the recipe generalizes to expression coordinates it was not fit on—yet the absolute value is well below 1. We read this as the central honest result of the paper: a real biological repolarization is *partly*, not fully, reachable by loss-of-function alone.

The signed decomposition (Fig. 2b) explains why. Only 39% of the target shift lies in the knockdown cone; 25% requires gain-of-function, and the remaining 35% aligns with neither. The activation certificate (Fig. 2c) makes the gain-of-function budget concrete, naming seven genes—LYAR, IKZF3, CRTAM, CBLB, GBP5, LAG3, and IRF8—that the T_H1 target requires to be raised but that no non-negative combination of knockdowns can deliver. These are not a null result but a set of falsifiable CRISPRa hypotheses: they specify exactly which activations a follow-up experiment should test.

The geometry recovers known biology as a positive control (Fig. 2d). GATA3, the master T_H2 transcription factor that must be removed to become T_H1 , is reached on the loss-of-function side at rank 155/6871 (97.7th percentile). TBX21, the master T_H1 factor that must instead be *raised*, is

correctly anti-aligned under knockdown at rank 6775/6871 (1.4th percentile)—the oracle refuses to knock down the gene the target needs induced. Both regulators are already reported by the source screen [10]; their recovery validates the operator and is not claimed as a discovery. The first genes entered by the greedy recipe (LAT2, APPBP2, RARA, ICOS) are likewise consistent with T-cell signaling and differentiation biology rather than novel nominations.




Figure 2: $T_H2 \rightarrow T_H1$ is partly reachable by knockdown. (a) Greedy reachability spectrum: reachable cosine versus recipe size k , far above the shuffled-target null at every k , with a knee at $k = 7$ (cosine 0.40) and a held-out cosine of 0.448 (null mean -0.005 , $z \approx 24$). (b) Signed decomposition of the target: 39% reachable by knockdown, 25% activation-required, 35% neither. (c) Activation certificate: seven genes the T_H1 target needs raised that no knockdown mix supplies—falsifiable CRISPRa hypotheses. (d) Positive control: the T_H2 driver GATA3 is reached (rank 155/6871) while the T_H1 driver TBX21 is correctly anti-aligned (rank 6775/6871).

3.2 Knockdown is never the majority modality across a 12-cell atlas

To test whether the partial-reachability finding is idiosyncratic to one transition, we ran the oracle across an atlas of 12 target cell-states: four transitions (toward T_H1 , toward T_H2 , toward younger, toward older) under three activation conditions (resting, stimulated 8 h, stimulated 48 h). Every one of the 12 targets is statistically reachable, with held-out z ranging from 4.7 to 45.0; none is a false positive of the cone’s dimensionality. But in no case does loss-of-function account for the

majority of the shift (Fig. 3a): the mean loss-of-function fraction is 0.34, and the knockdown bar never crosses the 50% line. Resting cells are the most reachable along all four axes—consistent with activation adding gain-of-function directions that knockdown cannot span—and the verdicts range from “partially reachable” (four resting transitions) to “weakly reachable” (the stimulated conditions). The general conclusion is robust: across transitions, cell states, and activation conditions, engineering by knockdown alone reaches a real but minority fraction of the desired change, and an honest method must report that fraction rather than a single reachable/not-reachable bit.

The atlas also exposes a calibration hazard that we report against ourselves (Fig. 3b). The in-sample reachable cosine (mean 0.580) systematically overstates the held-out-gene cosine (mean 0.355), by a mean gap of 0.225 (range 0.18–0.27) across all 12 transitions; an 85% without-replacement gene-panel bootstrap gives tight confidence intervals (e.g. toward T_{H1} resting: point 0.627, bootstrap mean 0.639, CI [0.633, 0.649]), confirming that the gap is a real optimism bias and not estimation noise. We therefore assign a confidence tier to each transition from its held-out z (2 high, 6 medium, 4 low) and quote held-out numbers everywhere. Finally, the closed-form anisotropy null lands almost exactly on the empirical shuffle null (Fig. 3c; Pearson $r = 0.995$ in-sample, 0.998 held-out), which is what makes the atlas-wide analysis affordable.

3.3 The verdict survives five reinforcement stress tests

A feasibility verdict is only trustworthy if it is not an artifact of a modelling or scoring choice. Five analyses each attack a different assumption; the headline conclusion—partly reachable by knockdown, with a named activation remainder—holds under all five, and two of them *tighten* the original claims rather than merely surviving. The analyses that carry a dedicated panel appear in Supplementary Figs. S5, S6, S8 and S9.

First, the non-negativity constraint earns its place. Relative to an unconstrained least-squares fit that is free to run knockdowns in reverse, the biological constraint costs only a mean of 0.018 in held-out cosine, yet the unconstrained fit buys that improvement with $\approx 3,537$ biologically unrealizable negative weights; the cone recovers a mean 95% (range 86–103%) of the unconstrained cosine and beats the single best generator in all 12 cells. The constraint is nearly free and is the sole source of the Farkas/KKT certificate.

Second, the modest absolute cosine is close to its modality-imposed ceiling. Because a knockdown-only cone cannot express gain-of-function directions, the attainable held-out cosine is bounded by $\sqrt{f_{LOF}}$; the headline 0.448 is 71% of that ceiling (atlas mean 61%, range 49–71%), and the residual gap is 22–26% gain-of-function-locked. The “modest” number is therefore most of what knockdown can achieve in principle, not a tuning shortfall.

Third, the recipes are additive-safe with margin. Every recommended knee recipe is additive-safe in all 12 cells (per-recipe reliability 0.92–0.96), and the combined-magnitude saturation cap binds only at $k \approx 28$ knockdowns versus a knee of $k \approx 4$ –5—a $5.6\times$ margin on the headline recipe.

Fourth, and most stringently, the verdict strengthens rather than collapses when the signal-dilution critique of Mejia et al. [14] is applied (Fig. S5). Reweighting the metric by each gene’s target magnitude ($w_g = |d_g|$) raises the Tier-2 $T_{H2} \rightarrow T_{H1}$ reachable cosine from 0.627 to 0.803 at rest (and every stimulated condition and every held-out Norman double likewise increases), while the held-out z stays far above the significance floor ($28.3 \rightarrow 14.1$ at rest). The dynamic-range fraction barely moves ($0.447 \rightarrow 0.499$) because the shuffled floor rises in step ($0.346 \rightarrow 0.619$); the verdict’s placement between floor and ceiling is stable, so it is not a dilution artifact (Fig. S6). The unweighted metric remains the default and reproduces the published numbers exactly.

Fifth, the certificate is reproducible but its interpretation must be stated carefully (Figs. S8–S9). Under a random gene-axis split the top certificate genes hold their rank (cross-half Spearman




Figure 3: Across a 12-cell atlas, knockdown is never the majority modality. (a) Signed LOF/GOF/neither decomposition for all 12 targets (four transitions $\times$ three activation states); the loss-of-function (green) bar never crosses the 50% majority line, mean LOF fraction 0.34, all held-out z between 4.7 and 45.0. (b) In-sample reachable cosine overstates the honest held-out cosine on every transition (points below $y = x$), mean gap 0.225; marker shade encodes the confidence tier. (c) The closed-form analytic null matches the empirical shuffle null (near-identity), justifying its use atlas-wide.

$\rho = 0.65$, half-versus-full 0.84, $z \approx 35$ against a shuffle null of ≈ 0 ; the full top-30 is retained), while the tail drifts. A retrospective cross-reference of the top 20 genes, however, qualifies the biological reading: the highest-demand genes are state-markers (LYAR, CRTAM, GBP5) and *negative* regulators (IKZF3/Aiolos, CBLB, LAG3)—genes whose *deletion*, not activation, boosts T-cell function—while the canonical T_H1 master transcription factors are present in the certificate but rank deep (IFNG #35, TBX21 #144, STAT4 #364, STAT1 #473). Of the 20, four are literature-SUPPORTED, seven PLAUSIBLE, and nine UNKNOWN, and 8/20 carry a strong (≥ 0.5) Open Targets genetic link to a T_H1 -driven autoimmune disease. We therefore read the certificate as a reproducible ranking of *unmet upward demand*, not as a validated list of activation targets: for the negative regulators in particular, the direction of a useful intervention may be opposite to naive activation. This is a tightening of the original “activation certificate” framing, and we carry it into the discussion.

Finally, applying the operator to an independent pair of cell types bounds what transfers, which

we develop next.

3.4 The operator transfers unchanged to a K562 CRISPRa screen

A feasibility operator is only useful if it is not bespoke to one screen. We applied the identical pipeline—same code, no retuned parameters—to the K562 combinatorial CRISPR activation screen of Norman et al. [13] (GSE133344), a different cell type, a different perturbation direction (activation rather than interference), and a different experimental design. Holding out the CEBPA master-regulator state as the target, the oracle reaches a cosine of 0.878 (residual norm 0.479) against a shuffled-target null with median 0.389 and 95th percentile 0.411 ($z \approx 37$ over 500 iterations), and a held-out-gene cosine of 0.856 ($z \approx 23.5$) (Fig. 4a). The recipe is led by CEBPE, the CEBPA paralog, with a knee at $k = 2$ (CEBPE + SPI1)—a biologically sensible minimal program. Where the target is not fully reached, the certificate again names an interpretable program: the CEBPA state requires raising a myeloid-differentiation module (MND4, HP, VSIG4, NCF1) orthogonal to what the activation atoms in the screen supply.

Crucially, the Norman screen contains *measured* double perturbations, which let us test the additivity assumption that underlies every multi-gene recipe rather than merely assert it (Fig. 4b). Across 126 measured doubles, the cosine between the true double perturbation and the sum of its two singles has median 0.71: additivity is a good but *bounded* approximation, not an identity. We characterized the departure directly. The deficit grows with the combined perturbation magnitude (Spearman +0.58), consistent with a scale-free saturation ceiling ($M^* = 13.9$ in units of the median single-effect norm, fit $R^2 = 0.57$); the intuitive hypothesis that collinear pairs should be most sub-additive was *refuted* (Spearman -0.16 , not significant). The practical consequence, which we carry into every recipe claim, is that multi-gene knockdown recipes are an extrapolation to be tested, not an exact prediction, and their reliability degrades at large combined magnitude.

To probe generalization more systematically than a single held-out state, we applied the operator to an independent pair of cell types—the genome-scale essential-gene Perturb-seq screens of Replogle et al. [15] in K562 and RPE1 cells (843 shared single-gene perturbations over 2,832 shared genes)—and asked, at three levels, what transfers across cell type (Supplementary Fig. S7). A single perturbation’s effect *direction* transfers moderately: the matched cross-type cosine has median 0.35 against a shuffled-gene null of ≈ 0 , rising to 0.64 along the common essential-stress axis. The reachability *verdict* transfers only as a graded quantity, not a binary flip—the residual norm is systematically worse under a cross-type basis than a within-type one ($0.61 \rightarrow 0.87$ for K562 targets, $0.37 \rightarrow 0.68$ for RPE1)—so a naive 100% verdict-agreement rate is an artifact of an over-expressive 843-atom cone and should be read as a residual, not a label. The minimal *recipe* does *not* transfer: the same-target Jaccard overlap across cell types is only 0.11, collapsing to the random-selection null (0.05) under a cross-type basis. The prescription is therefore bounded: the *direction* of a cell-state shift is portable across cell types, but the *specific minimal recipe* is cell-type-specific and must be re-derived per system. This is a cross-cell-type generalization probe on essential-gene screens, not a second CD4⁺ T-cell polarization screen.

3.5 No prior method occupies the measured-effects, verdict-plus-certificate corner

We placed the oracle against a structured survey of 92 methods spanning 2011–2026 (Section 4), scored on two axes: how strongly each is grounded in measured effect vectors, and whether it emits a feasibility verdict with a certificate (Fig. 5a). The two capabilities are held by *disjoint* sets of prior methods. The 14 methods that operate directly on measured effects are all forward-only predictors




Figure 4: Transfer to Norman 2019 K562 CRISPRa and the additivity guardrail. (a) The unchanged operator reaches the held-out CEBPA state at cosine 0.878 (full-fit), far above the shuffled-target null ($z \approx 37$; held-out-gene $z \approx 23.5$), with the recipe led by the CEBPA paralog CEBPE. (c) The infeasible remainder’s certificate names a myeloid program (MNDA, HP, VSIG4, NCF1). (b) Across 126 measured double perturbations, the cosine of the measured double to the sum of its singles has median 0.71: additivity is a bounded approximation, and the “epistasis gap” to perfect additivity ($= 1$) is real.

or rankers; none returns a feasibility verdict or a data-grounded certificate. The 13–15 methods carrying any verdict- or certificate-like output are all model-theoretic—their guarantee is a property of a hand-built, fitted, or inferred model, and they span several families (structural and Boolean network control, control-from-data, an inferred single-cell GRN, a causally inspired predictor) rather than one. The intersection of “grounded in measured effects” and “emits a verdict plus certificate” is empty for all 91 prior methods; the reachability oracle is the sole occupant of that corner. The family-level capability matrix (Supplementary Fig. S1) shows the same structure per capability column: measured-effect grounding and feasibility output never co-occur in any prior family.

We state the scope of this claim precisely. It is a claim about the *combination* of two capabilities on a structured survey, not a claim that any individual capability is unprecedented; network control has reasoned about reachability for over a decade, and many methods use measured data. The novelty is the pairing—a feasibility decision computed directly from measured effects, with a certificate—and the disjointness is the evidence for it.




Figure 5: Landscape positioning and modality triage. (a) Each of 92 surveyed methods (2011–2026) placed by measured-effect grounding (x) and feasibility-verdict-plus-certificate output (y). Measured-grounded methods (blue) are forward-only; verdict/certificate methods (pink) are model-theoretic; the top-right corner is empty across all 91 prior methods and occupied only by this work. (b) The verdict translates into a GO/REDIRECT/STOP triage over 102 real target nominations (45/102 = 44% hard-to-drug, 10 with a clinical-grade drug); all nominations are falsifiable wet-lab hypotheses, not validated targets.

3.6 From verdict to experiment: a GO/REDIRECT/STOP triage

The point of a feasibility verdict is that it changes what experiment gets run. We applied the oracle to 102 real target nominations drawn from the T-cell transitions and translated each verdict into one of three actions (Fig. 5b): **GO**—the target is reachable, so hand the minimal ranked knockdown set to an arrayed CRISPRi screen; **REDIRECT**—the target is activation-required, so switch to a CRISPRa/agonist arm, with the certificate naming the genes to test first; and **STOP**—the target is provably outside the cone, so do not spend a knockdown arm on it. This matters because 45/102 (44%) of the nominated targets are hard-to-drug and only 10 carry a clinical-grade drug, so misallocating screen arms is expensive. Across transitions, 36–49% (mean $\approx$ 42%) of each transition’s addressable shift needs induction and is therefore unreachable by a knockdown-only screen—exactly the fraction a CRISPRi-only campaign would waste effort on without a feasibility check.

The triage surfaces concrete, and honestly caveated, examples. Nine targets are green-lit as directly reachable and druggable (including JAK2, ICOS, MAPK14, CD3D). Three are tractable but untried and flagged for a redirected arm (IL7R, which carries the highest genetic support; ZAP70; TET2). One, IRF1, is a top T_H2 knockdown node with strong immune-disease genetic support but no conventional small-molecule or antibody handle (degrader-only), a required-but-undruggable case where the verdict argues for a new modality rather than a conventional campaign. We emphasize that every one of these is a prioritization hypothesis for the bench, generated by geometry over

one screen; none is a validated target, and the oracle’s role is to triage where experimental effort is spent, not to replace the experiment.

4 Related work

Cell-state engineering poses an inverse-design question: given a starting state and a desired target state, which interventions move the cell there, and is the move even possible? This is the question behind directed differentiation, transdifferentiation, and T-cell engineering, and it is the question of pharmaceutical target selection. Getting it wrong is the dominant cost of drug development: only about one in ten clinical programmes reaches approval, and failure is paid for late [1]. Attrition concentrates at Phase II, and roughly half of Phase II/III failures are for lack of efficacy—the wrong target or mechanism rather than safety [3]. The two interventions with published evidence of improving those odds both point the same way: human genetic support raises the odds of approval about 2.6-fold [1, 2], and *measured* rather than *inferred* target effects are the direct antidote to the 11–25% preclinical reproducibility documented by the Amgen and Bayer audits [4, 5]. A method for this problem should therefore (i) work from measured effects in the relevant cell type, (ii) answer achievability rather than merely rank candidates, and (iii) say so honestly when a goal is not achievable.

We surveyed 91 prior methods spanning 2011–2026 (92 including this work) across control-theoretic network control, gene-regulatory-network (GRN) inference with in-silico perturbation, deep perturbation-response prediction, single-cell foundation models, and the combinatorial-screen, optimal-transport, virtual-cell, and benchmark literature that borders them. Positioning them on two axes—*what question is answered* (forward prediction versus feasibility) and *what the answer is grounded in* (an inferred or hand-built model versus directly measured effects)—exposes a gap that is exact rather than rhetorical. The methods grounded in measured effects answer only forward questions, and the methods that reason about feasibility do so only inside a model.

Model-theoretic feasibility. The most mature theory for the intervention question is network control. Structural and target controllability cast steering an assumed linear system as a matching problem over network topology [16, 17], and determining-node / control-kernel methods identify driver sets from structure [18, 19]. Boolean attractor control computes node interventions that switch a logical network between attractors [20–23], a line that remains active [24–28]. A data-fitted strand fits a dynamical model to omics and then applies control theory [29–31], and network-influence predictors rank transcription-factor sets for a conversion [32, 33]. These are the methods that carry a controllability or control-strategy guarantee—but the guarantee is a property of a hand-built, fitted, or inferred model, never of measured effects, and the control set is only as trustworthy as the network put in. The same dependency reappears whenever feasibility-style reasoning is imported onto a reconstructed single-cell GRN [34] or run in a learned latent space [35]. The set of prior methods that emit any verdict- or certificate-like output is therefore not one family: it spans structural and Boolean network control, control-from-data, an inferred single-cell GRN, and one causally inspired perturbation predictor. What unites them is the assumed model, not the family.

Measured-grounded forward prediction. The single-cell community answers the same biological question by running an inferred network forward. CellOracle infers a GRN and simulates a transcription-factor knockout into predicted identity transitions [36]; related tools infer regulons, enhancer networks, or dynamic GRNs [37–41], with BEELINE quantifying how method-dependent GRN inference is [42]. Deep perturbation-response models are genuinely measured-grounded—they

learn from real Perturb-seq data—but are resolutely forward: GEARS predicts outcomes of unseen multi-gene perturbations [43], and compositional and disentanglement models predict responses to given perturbations [44–47]. Single-cell foundation models [48–50] and the 2025–2026 virtual-cell surge—perturbation transformers, atlas-scale models and challenges, and giga-scale perturbation atlases [51–54]—extend the same forward map. Optimal-transport methods describe how cells *do* move rather than which interventions *would* move them [55, 56], and experimental-design methods choose the next screen [57]. Every one maps perturbation to predicted response: a predictor always returns *some* answer and structurally cannot certify that a target is unreachable.

Why a feasibility verdict, not another predictor. A cluster of 2025–2026 benchmarks reports that current perturbation-prediction models, foundation models included, often fail to beat simple mean or linear baselines on genuinely unseen perturbations [6–9, 58–65]. This is the field’s own critique, and it sharpens the case for a different output: a feasibility verdict computed directly from measured effects does not extrapolate to unseen perturbations at all, and it comes with a checkable certificate and a shuffled-target null. The claim is deliberately weaker, and therefore more defensible, than predicting an unseen response.

Our position. Across all 91 prior methods (ours excluded), 14 operate directly on measured effects—and all 14 only predict or rank; *none* returns a feasibility verdict, emits a data-grounded certificate, or proves infeasibility, and the 13–15 partial verdict/certificate marks are all model-theoretic. The two capabilities this work combines—measured grounding and a feasibility verdict with a certificate—are held by *disjoint* sets of prior methods, and their intersection is empty. We compose the measured genome-scale CRISPRi effect vectors from a Perturb-seq screen in primary human CD4⁺ T cells [10] directly, with no inferred network. Because knockdown is loss-of-function and combinations are non-negative, reachability becomes membership in a convex cone, solved by non-negative least squares [11]: “reachable” returns the mixing weights (a minimal ranked recipe), and “outside the cone” is a theorem whose separating hyperplane—a Farkas/KKT certificate [12]—names the genes that would have to be *activated* instead, a falsifiable CRISPRa hypothesis rather than a null result. We validate by held-out-gene generalisation against a shuffled-target null (Th2→Th1 held-out cosine 0.448, null $z \approx 24$; KKT residual $\approx 10^{-11}$), and the operator transfers unchanged to the Norman K562 combinatorial CRISPR screen [13] (held-out CEBPA cosine 0.878), whose measured double-perturbations show that the additivity assumption is a bounded approximation (median cosine 0.71, sum-of-singles versus measured) rather than exact. We claim novelty for this measured-effect, certificate-carrying reachability *decision*, not for the T-cell or CEBPA regulator lists, which the source screens [10, 13] already report and which we use only as positive controls; activation hypotheses require a separate CRISPRa arm [66], and the modest absolute cosine is the honest ceiling of loss-of-function alone.

5 Discussion

We have described a feasibility oracle that treats a genome-scale CRISPRi screen as a dictionary of measured effect vectors and decides, by convex geometry, whether a target cell-state shift can be reached by a non-negative combination of knockdowns. The decision is constructive on both sides: a reachable target returns a minimal ranked recipe, and an unreachable one returns a Farkas/KKT certificate that names the genes requiring activation. The value of the framework is not a new list of T-cell regulators—those are already in the source screens and we use them only as positive

controls—but a new *kind of output*: a falsifiable feasibility verdict computed directly from measured effects, which no method in a 92-entry survey provides.

5.1 Interpretation

Three findings are worth restating for their scientific rather than methodological content. First, a real polarization shift ($T_{H2} \rightarrow T_{H1}$) is only *partly* reachable by knockdown (held-out cosine 0.448; 39% of the shift in the loss-of-function cone), and this pattern holds across a 12-cell atlas where knockdown is never the majority modality (mean loss-of-function fraction 0.34). The implication for practice is direct: a knockdown-only campaign, however well designed, addresses a minority of most cell-state transitions, and the remaining majority is a mix of required activation and an irreducible remainder. This gap is not a tuning shortfall: the headline held-out cosine is 71% of the modality-imposed ceiling that a knockdown-only cone can attain in principle (atlas mean 61%), so most of what loss-of-function can supply is already captured. Second, the certificate turns this limitation into a plan—it names, reproducibly, which genes carry the unmet upward demand—so an infeasibility result redirects rather than terminates an experiment. We are deliberately careful about how that list is read: the highest-demand genes are state-markers and *negative* regulators rather than the canonical master transcription factors (which are present but rank deep), so the certificate is a reproducible statement of *where activation is required*, not a validated list of activation targets, and for some entries the useful intervention may be opposite in sign to naive activation. Third, the operator transfers unchanged to screens in other cell types, but transfer is graded: the *direction* of a cell-state shift is portable across cell types while the *specific minimal recipe* is not, so the framework is a property of the geometry whose quantitative recipe must still be re-derived per system.

5.2 Limitations

The rigor of a feasibility claim rests on being explicit about what it does not establish. We group the principal limitations by severity; a detailed, actionable reinforcement plan accompanies this manuscript as a standalone document.

The additivity assumption is load-bearing and only bounded. Every multi-gene recipe assumes that co-perturbation effects add. On the 126 measured double perturbations available (Norman K562), the measured double aligns with the sum of its singles at median cosine 0.71—good, but not exact, and degrading with combined perturbation magnitude (Spearman +0.58 deficit-versus-magnitude; a scale-free saturation ceiling $M^* = 13.9$). The intuitive hypothesis that collinear perturbations should be most sub-additive was refuted (Spearman -0.16 , not significant), so the departure is a magnitude effect, not a redundancy effect. This is the single most important caveat: recipes of more than a few knockdowns, or knockdowns of individually large effect, are extrapolations whose error is bounded but real, and the CD4⁺ T-cell dictionary itself contains no measured doubles with which to check additivity in-domain.

Loss-of-function has a structural ceiling. The modest absolute reachable cosine is not a tuning failure but a property of the modality: a non-negative combination of knockdowns cannot express directions that require increased expression, so any target with a substantial gain-of-function component is unreachable by construction. The knockdown-only ceiling equals $\sqrt{f_{LOF}}$ and matches the in-sample cone cosine to 10^{-4} ; against it, the headline 0.448 is 71% of what is attainable in principle (atlas mean 61%, range 49–71%), with the residual 22–26% gain-of-function-locked. This is why we report a signed decomposition rather than a single reachability score, and why the honest

headline number is well below one *yet* close to its own ceiling—a distinction we make explicitly so the number is not read as a failure of fit.

One perturbation system, few donors, collapsed across donors. The primary results derive from a single primary-cell system (CD4⁺ T cells) with four donors, and the effect vectors available to us are collapsed across donors, so we cannot perform a true leave-one-donor-out generalization test on the primary dictionary. External validity across genotypes, tissues, and laboratories is therefore asserted by the cross-dataset transfer to K562, not yet established within the primary system. The transcription-factor-versus-marker split in the positive-control analysis is likewise a post-hoc, exploratory annotation rather than a pre-registered test.

Statistical nulls answer narrow questions. The shuffled-target null establishes that a target is more reachable than a structure-matched random direction; it does not establish that the recipe is biologically correct. The held-out-gene analysis establishes generalization across expression coordinates; it does not establish generalization across cells or conditions. And the in-sample cosine systematically overstates the held-out value (mean gap 0.225 across the atlas), which is why every headline number in this paper is the held-out one. A separate random-perturbation diagnostic is negative by construction and is used only as an internal control, not as evidence of performance.

The certificate is reproducible but not yet a validated target list. Under a random gene-axis split the top certificate genes hold their rank (cross-half Spearman $\rho \approx 0.65$, half-versus-full ≈ 0.84 , $z \approx 35$ against a shuffle null), so the ranking is a stable property of the data and not of a particular fit. Its *biological* interpretation is more guarded. A retrospective cross-reference of the top 20 genes finds them dominated by state-markers and negative regulators, with the canonical T_H1 master transcription factors present but ranked deep; four are literature-supported, seven plausible, nine unknown, and 8/20 carry a strong genetic link to a T_H1-driven autoimmune disease. We therefore read the certificate as a reproducible ranking of unmet upward demand, and we flag that for the negative regulators an “activate this gene” reading may be directionally wrong—the correct intervention could be further repression. This cross-reference is exploratory and post-hoc, and it tempers rather than confirms the earlier “immunologically credible activation certificate” framing.

Transfer is graded, and the recipe does not port across cell types. Applying the operator to an independent pair of cell types (K562 and RPE1 essential-gene screens) shows that transfer is a spectrum, not a switch. Single-perturbation effect *direction* transfers moderately (matched cross-type cosine median 0.35 versus a ≈ 0 shuffled-gene null), and the reachability verdict transfers as a graded residual rather than a binary flip; but the minimal *recipe* does not transfer (same-target Jaccard 0.11, collapsing to the random null under a cross-type basis). The apparent $\approx 100\%$ verdict-agreement rate across cell types is an artifact of an over-expressive 843-atom cone and must be read as a residual, not a label. The practical bound is that a recipe derived in one cell type is a hypothesis, not a prescription, in another.

The verdict is robust to the choice of evaluation metric. Uniform-gene metrics can in principle dilute a real signal by averaging over unchanged genes. Reweighting every metric by target magnitude ($w_g = |d_g|$)—the correction called for by Mejia et al. [14]—*strengthens* rather than weakens the headline verdict (Tier-2 rest reachable cosine $0.627 \rightarrow 0.803$; held-out z remains far above the significance floor), and the calibrated dynamic-range placement is stable because the

shuffled floor rises in step. We report the unweighted metric as the default because it reproduces the published numbers exactly, but the conclusion does not depend on that choice.

5.3 Outlook

The most valuable next experiments follow directly from the limitations. Measuring a modest panel of double knockdowns in the primary CD4⁺ system would replace the borrowed K562 additivity estimate with an in-domain one and calibrate the recipe-size ceiling. Running one CRISPRa arm against a certificate’s named genes would provide the first direct test of an infeasibility prediction—the framework’s most distinctive and least-validated claim. And extending the dictionary to additional cell types and donors would convert the transfer result from an existence proof into a generalization curve. In each case the framework’s output is what makes the experiment efficient: it says not only what to test but where testing is unnecessary. We therefore position the oracle as an experiment-triage instrument—a way to decide where a screen should spend its arms and where it should stop—rather than a target-validation engine, and we release it in that spirit.

References

- [1] Eric Vallabh Minikel, Jeffery L. Painter, Coco Chengliang Dong, and Matthew R. Nelson. Refining the impact of genetic evidence on clinical success. *Nature*, 629(8012):624–629, 2024. doi: 10.1038/s41586-024-07316-0.
- [2] Matthew R. Nelson, Hannah Tipney, Jeffery L. Painter, Judong Shen, Paola Nicoletti, Yufeng Shen, Aris Floratos, Pak C. Sham, Mulin Jun Li, Junwen Wang, Lon R. Cardon, John C. Whittaker, and Philippe Sanséau. The support of human genetic evidence for approved drug indications. *Nature Genetics*, 47(8):856–860, 2015. doi: 10.1038/ng.3314.
- [3] Richard K. Harrison. Phase II and phase III failures: 2013-2015. *Nature Reviews Drug Discovery*, 15(12):817–818, 2016. doi: 10.1038/nrd.2016.184.
- [4] C. Glenn Begley and Lee M. Ellis. Raise standards for preclinical cancer research. *Nature*, 483(7391):531–533, 2012. doi: 10.1038/483531a.
- [5] Florian Prinz, Thomas Schlange, and Khusru Asadullah. Believe it or not: how much can we rely on published data on potential drug targets? *Nature Reviews Drug Discovery*, 10(9):712–712, 2011. doi: 10.1038/nrd3439-c1.
- [6] Constantin Ahlmann-Eltze, Wolfgang Huber, and Simon Anders. Deep-learning-based gene perturbation effect prediction does not yet outperform simple linear baselines. *Nature Methods*, 22(8):1657–1661, 2025. doi: 10.1038/s41592-025-02772-6.
- [7] Daniel R. Wong, Abby S. Hill, and Rob Moccia. Simple controls exceed best deep learning algorithms and reveal foundation model effectiveness for predicting genetic perturbations. *Bioinformatics*, 41(6), 2025. doi: 10.1093/bioinformatics/btaf317.
- [8] Eric Kernfeld, Yunxiao Yang, Joshua S. Weinstock, Alexis Battle, and Patrick Cahan. A comparison of computational methods for expression forecasting. *Genome biology*, 26(1):388–388, 2025. doi: 10.1186/s13059-025-03840-y.

- [9] Srijan Atti and Shankar Subramaniam. Fundamental Limitations of Foundation Models in Single-Cell Transcriptomics. *bioRxiv (Cold Spring Harbor Laboratory)*, 2025. doi: 10.1101/2025.06.26.661767. Preprint.
- [10] Ronghui Zhu, Emma Dann, Jun Yan, Justine Reyes Retana, Ryunosuke Goto, Reese C Guitche, Lillian Petersen, Mineto Ota, Jonathan K. Pritchard, and Alexander Marson. Genome-scale perturb-seq in primary human CD4+ T cells maps context-specific regulators of T cell programs and human immune traits. *bioRxiv*, 2025. doi: 10.64898/2025.12.23.696273. Preprint. bioRxiv DOI on the openRxiv 10.64898 prefix (post-2025 migration); resolves via OpenAlex and the bioRxiv API.
- [11] C. L. Lawson and Richard Hanson. *Solving Least Squares Problems*. Classics in Applied Mathematics. Society for Industrial and Applied Mathematics (SIAM), 1974. doi: 10.1137/1.9781611971217. Originally Prentice-Hall, 1974; SIAM Classics reissue 1995.
- [12] Julius Farkas. Theorie der einfachen Ungleichungen. *Journal für die reine und angewandte Mathematik*, 1902(124):1–27, 1902. doi: 10.1515/crll.1902.124.1.
- [13] Thomas M. Norman, Max A. Horlbeck, Joseph M. Replogle, Alex Y. Ge, Albert Xu, Marco Jost, Luke A. Gilbert, and Jonathan S. Weissman. Exploring genetic interaction manifolds constructed from rich single-cell phenotypes. *Science*, 365(6455):786–793, 2019. doi: 10.1126/science.aax4438.
- [14] Gabriel Mateo Mejia, Henry E. Miller, Francis J. A. Leblanc, Bo Wang, Brendan Swain, and Lucas Paulo de Lima Camillo. Needles in the Haystack: Addressing Signal Dilution Improves scRNA-seq Perturbation Response Modeling and Evaluation. In *Proceedings of the 43rd International Conference on Machine Learning (ICML)*, 2026. ICML 2026 poster (OpenReview XsrXLPxBJw). Expanded preprint: Miller et al., bioRxiv 2025.10.20.683304, doi:10.1101/2025.10.20.683304, “Deep Learning-Based Genetic Perturbation Models Do Outperform Uninformative Baselines on Well-Calibrated Metrics”.
- [15] Joseph M. Replogle, Reuben A. Saunders, Angela N. Pogson, Jeffrey A. Hussmann, Alexander Lenail, Alina Guna, Lauren Mascibroda, Eric J. Wagner, Karen Adelman, Gila Lithwick-Yanai, Nika Iremadze, Florian Oberstrass, Doron Lipson, Jessica L. Bonnar, Marco Jost, Thomas M. Norman, and Jonathan S. Weissman. Mapping information-rich genotype-phenotype landscapes with genome-scale Perturb-seq. *Cell*, 185(14):2559–2575.e28, 2022. doi: 10.1016/j.cell.2022.05.013.
- [16] Yang-Yu Liu, Jean-Jacques Slotine, and Albert-Ászló Barabási. Controllability of complex networks. *Nature*, 473(7346):167–173, 2011. doi: 10.1038/nature10011.
- [17] Jianxi Gao, Yang-Yu Liu, Raissa M. D’Souza, and Albert-Ászló Barabási. Target control of complex networks. *Nature Communications*, 5(1):5415–5415, 2014. doi: 10.1038/ncomms6415.
- [18] Atsushi Mochizuki, Bernold Fiedler, Gen Kurosawa, and Daisuke Saito. Dynamics and control at feedback vertex sets. II: A faithful monitor to determine the diversity of molecular activities in regulatory networks. *Journal of Theoretical Biology*, 335:130–146, 2013. doi: 10.1016/j.jtbi.2013.06.009.
- [19] Junil Kim, Sang-Min Park, and Kwang-Hyun Cho. Discovery of a kernel for controlling biomolecular regulatory networks. *Scientific Reports*, 3(1):2223–2223, 2013. doi: 10.1038/srep02223.

- [20] Jorge Gómez Tejeda Zañudo and Réka Albert. Cell Fate Reprogramming by Control of Intracellular Network Dynamics. *PLoS Computational Biology*, 11(4):e1004193–e1004193, 2015. doi: 10.1371/journal.pcbi.1004193.
- [21] Jorge Gómez Tejeda Zañudo, Gang Yang, and Réka Albert. Structure-based control of complex networks with nonlinear dynamics. *Proceedings of the National Academy of Sciences*, 114(28):7234–7239, 2017. doi: 10.1073/pnas.1617387114.
- [22] Jae Il Joo, Joseph Zhou, Sui Huang, and Kwang-Hyun Cho. Determining Relative Dynamic Stability of Cell States Using Boolean Network Model. *Scientific Reports*, 8(1):12077–12077, 2018. doi: 10.1038/s41598-018-30544-0.
- [23] Lauren Marazzi, Milan Shah, Shreedula Balakrishnan, Ananya Patil, and Paola Vera-Licona. NETISCE: a network-based tool for cell fate reprogramming. *npj Systems Biology and Applications*, 8(1):21–21, 2022. doi: 10.1038/s41540-022-00231-y.
- [24] Laura Cifuentes-Fontanals, Elisa Tonello, and Heike Siebert. Node and edge control strategy identification via trap spaces in Boolean networks. *BMC Bioinformatics*, 24(S1):494–494, 2025. doi: 10.1186/s12859-025-06135-y.
- [25] John Metzcar, Katie Pletz, Heber L. Rocha, and Jordan C. Rozum. Translating and evaluating single-cell Boolean network interventions in the multiscale setting. *arXiv (Cornell University)*, 2025. doi: 10.48550/arxiv.2501.16052. Preprint.
- [26] David Murrugarra, Alan Veliz-Cuba, Elena Dimitrova, Claus Kadelka, Matthew T. Wheeler, and Reinhard Laubenbacher. Modular Control of Boolean Network Models. *Bulletin of Mathematical Biology*, 87(7):91–91, 2025. doi: 10.1007/s11538-025-01471-9.
- [27] Van-Giang Trinh, Kyu Hyong Park, Samuel Pastva, and Jordan C. Rozum. Mapping the attractor landscape of Boolean networks with biobalm. *Bioinformatics*, 41(5), 2025. doi: 10.1093/bioinformatics/btaf280.
- [28] Domitilla Del Vecchio. Control systems for synthetic biology and a case-study in cell fate reprogramming. *Open MIND*, 2026. doi: 10.48550/arxiv.2601.20135. Preprint.
- [29] Scott Ronquist, Geoff Patterson, Lindsey A. Muir, Stephen Lindsly, Haiming Chen, Markus Brown, Max S. Wicha, Anthony M. Bloch, Roger W. Brockett, and Indika Rajapakse. Algorithm for cellular reprogramming. *Proceedings of the National Academy of Sciences*, 114(45):11832–11837, 2017. doi: 10.1073/pnas.1712350114.
- [30] Oleksii S. Rukhlenko, Melinda Halász, Nora Rauch, Vadim Zhernovkov, Thomas L. Prince, Kieran Wynne, Stephanie Maher, Eugene Kashdan, Kenneth Macleod, Neil O. Carragher, Walter Kölch, and Boris N. Kholodenko. Control of cell state transitions. *Nature*, 609(7929):975–985, 2022. doi: 10.1038/s41586-022-05194-y.
- [31] Thomas P. Wytock and Adilson E. Motter. Cell reprogramming design by transfer learning of functional transcriptional networks. *Proceedings of the National Academy of Sciences*, 2024. doi: 10.1073/pnas.2312942121. Preprint.
- [32] Owen J. L. Rackham, Jaber Firas, Hai Fang, Matt E. Oates, Melissa L. Holmes, Anja S. Knaupp, Harukazu Suzuki, Christian M. Nefzger, Carsten O. Daub, Jay W. Shin, Enrico Petretto, Alistair R. R. Forrest, Yoshihide Hayashizaki, José M. Polo, and Julian Gough.

- A predictive computational framework for direct reprogramming between human cell types. *Nature Genetics*, 48(3):331–335, 2016. doi: 10.1038/ng.3487.
- [33] Patrick Cahan, Hu Li, Samantha A. Morris, Edroaldo Lummertz da Rocha, George Q. Daley, and James J. Collins. CellNet: Network Biology Applied to Stem Cell Engineering. *Cell*, 158(4):903–915, 2014. doi: 10.1016/j.cell.2014.07.020.
 - [34] Peizhuo Wang, Wen Xiao, Han Li, Lang Peng, Shuya Li, Yipin Lei, Hantao Shu, Lin Gao, Dan Zhao, and Jianyang Zeng. Deciphering driver regulators of cell fate decisions from single-cell transcriptomics data with CEFCON. *Nature Communications*, 14(1):8459–8459, 2023. doi: 10.1038/s41467-023-44103-3.
 - [35] Young-Hyun Han, Hyunjin Kim, Chun-Kyung Lee, and Kwang-Hyun Cho. Identifying an optimal perturbation to induce a desired cell state by generative deep learning. *Cell Systems*, 16(10):101405–101405, 2025. doi: 10.1016/j.cels.2025.101405.
 - [36] Kenji Kamimoto, Blerta Stringa, Christy M. Hoffmann, Kunal Jindal, Lilianna Solnica-Krezel, and Samantha A. Morris. Dissecting cell identity via network inference and in silico gene perturbation. *Nature*, 614(7949):742–751, 2023. doi: 10.1038/s41586-022-05688-9.
 - [37] Sara Aibar, Carmen Bravo González-Blas, Thomas Moerman, Vân Anh Huynh-Thu, Hana Imrichová, Gert Hulselmans, Florian Rambow, Jean-Christophe Marine, Pierre Geurts, Jan Aerts, Joost van den Oord, Zeynep Kalender Atak, Jasper Wouters, and Stein Aerts. SCENIC: single-cell regulatory network inference and clustering. *Nature Methods*, 14(11):1083–1086, 2017. doi: 10.1038/nmeth.4463.
 - [38] Quan Xu, Georgios Georgiou, Siebren Frölich, Maarten van der Sande, Gert Jan C. Veenstra, Huiqing Zhou, and Simon J. van Heeringen. ANANSE: an enhancer network-based computational approach for predicting key transcription factors in cell fate determination. *Nucleic Acids Research*, 49(14):7966–7985, 2021. doi: 10.1093/nar/gkab598.
 - [39] Carmen Bravo González-Blas, Seppe De Winter, Gert Hulselmans, Nikolai Hecker, Irina Matetovici, Valerie Christiaens, Suresh Poovathingal, Jasper Wouters, Sara Aibar, and Stein Aerts. SCENIC+: single-cell multiomic inference of enhancers and gene regulatory networks. *Nature Methods*, 20(9):1355–1367, 2023. doi: 10.1038/s41592-023-01938-4.
 - [40] Lingfei Wang, Nikolaos Trasanidis, Ting Wu, Guanlan Dong, M. Hu, Daniel E. Bauer, and Luca Pinello. Dictys: dynamic gene regulatory network dissects developmental continuum with single-cell multiomics. *Nature Methods*, 20(9):1368–1378, 2023. doi: 10.1038/s41592-023-01971-3.
 - [41] Daniel Osorio, Yan Zhong, Guanxun Li, Qian Xu, Yongjian Yang, Yanan Tian, Robert S. Chapkin, Jianhua Z. Huang, and James J. Cai. scTenifoldKnk: An efficient virtual knockout tool for gene function predictions via single-cell gene regulatory network perturbation. *Patterns*, 3(3):100434–100434, 2022. doi: 10.1016/j.patter.2022.100434.
 - [42] Aditya Pratapa, Amogh P. Jalihal, Jeffrey Law, Aditya Bharadwaj, and T. M. Murali. Benchmarking algorithms for gene regulatory network inference from single-cell transcriptomic data. *Nature Methods*, 17(2):147–154, 2020. doi: 10.1038/s41592-019-0690-6.
 - [43] Yusuf Roohani, Kexin Huang, and Jure Leskovec. Predicting transcriptional outcomes of novel multigene perturbations with GEARS. *Nature Biotechnology*, 42(6):927–935, 2023. doi: 10.1038/s41587-023-01905-6.

- [44] Mohammad Lotfollahi, F. Alexander Wolf, and Fabian J. Theis. scGen predicts single-cell perturbation responses. *Nature Methods*, 16(8):715–721, 2019. doi: 10.1038/s41592-019-0494-8.
- [45] Mohammad Lotfollahi, Anna Klimovskaia Susmelj, Carlo De Donno, Leon Hetzel, Yuge Ji, Ignacio L. Ibarra, Sanjay Srivatsan, Mohsen Naghipourfar, Riza M. Daza, Beth Martin, Jay Shendure, José L. McFaline-Figueroa, Pierre Boyeau, F. Alexander Wolf, Nafissa Yakubova, Stephan Günemann, Cole Trapnell, David López-Paz, and Fabian J. Theis. Predicting cellular responses to complex perturbations in high-throughput screens. *Molecular Systems Biology*, 19(6):e11517–e11517, 2023. doi: 10.15252/msb.202211517.
- [46] Leon Hetzel, Simon Böhm, Niki Kilbertus, Stephan Günemann, Mohammad Nader Lotfollahi, and Fabian J. Theis. Predicting Cellular Responses to Novel Drug Perturbations at a Single-Cell Resolution. *NeurIPS 2022 (arXiv:2204.13545)*, 2022. doi: 10.48550/arXiv.2204.13545. Preprint.
- [47] Zoe Piran, Niv Cohen, Yedid Hoshen, and Mor Nitzan. Disentanglement of single-cell data with biolord. *Nature Biotechnology*, 42(11):1678–1683, 2024. doi: 10.1038/s41587-023-02079-x.
- [48] Christina V. Theodoris, Ling Xiao, Anant Chopra, Mark Chaffin, Zeina R. Al Sayed, Matthew C. Hill, Helene Mantineo, Elizabeth M. Brydon, Zexian Zeng, X. Shirley Liu, and Patrick T. Ellinor. Transfer learning enables predictions in network biology. *Nature*, 618(7965):616–624, 2023. doi: 10.1038/s41586-023-06139-9.
- [49] Haotian Cui, Chloe Wang, Hassaan Maan, Kuan Pang, Fengning Luo, Nan Duan, and Bo Wang. scGPT: toward building a foundation model for single-cell multi-omics using generative AI. *Nature Methods*, 21(8):1470–1480, 2024. doi: 10.1038/s41592-024-02201-0.
- [50] Minsheng Hao, Jing Gong, Xin Zeng, Chi-Ming Liu, Yucheng Guo, Xingyi Cheng, Taifeng Wang, Jianzhu Ma, Xuegong Zhang, and Le Song. Large-scale foundation model on single-cell transcriptomics. *Nature Methods*, 21(8):1481–1491, 2024. doi: 10.1038/s41592-024-02305-7.
- [51] Abhinav Adduri, Dhruv Gautam, Beatrice Bevilacqua, Alishba Imran, Rohan Shah, Mohsen Naghipourfar, Noam Teyssier, Rajesh Ilango, Sanjay Nagaraj, Mingze Dong, Chiara Ricci-Tam, Christopher Carpenter, Vishvak Subramanyam, Aidan Winters, Sravya Tirukkovular, Jeremy A. Sullivan, Brian S. Plosky, Basak Eraslan, Nicholas D. Youngblut, Jure Leskovec, Luke A. Gilbert, Silvana Konermann, Patrick D. Hsu, Alexander Dobin, Dave P. Burke, Hani Goodarzi, and Yusuf Roohani. Predicting cellular responses to perturbation across diverse contexts with State. *bioRxiv (Cold Spring Harbor Laboratory)*, 2025. doi: 10.1101/2025.06.26.661135. Preprint.
- [52] Yusuf Roohani, Tony J. Hua, Po-Yuan Tung, Lexi R. Bounds, Feiqiao Brian Yu, Alexander Dobin, Noam Teyssier, Abhinav Adduri, Alden Woodrow, Brian S. Plosky, Reshma Mehta, Benjamin Hsu, Jeremy A. Sullivan, Chiara Ricci-Tam, Nianzhen Li, Julia Kazaks, Luke A. Gilbert, Silvana Konermann, Patrick D. Hsu, Hani Goodarzi, and Dave P. Burke. Virtual Cell Challenge: Toward a Turing test for the virtual cell. *Cell*, 188(13):3370–3374, 2025. doi: 10.1016/j.cell.2025.06.008.
- [53] Jesse Zhang, Airol A Ubas, Richard de Borja, Valentine Svensson, Nicole Thomas, Neha Thakar, Ian Lai, Aidan Winters, Umair Khan, Matthew G. Jones, John D. Thompson, Vuong Tran, Joseph Pangallo, Efthymia Papalexi, Ajay Sapre, Hoai Nguyen, Oliver Sanderson, Maria Nigos, Olivia Kaplan, Sarah Schroeder, Bryan Hariadi, Simone Marrujo, Crina Curca Alec Salvino, Guillermo Gallareta Olivares, Ryan Koehler, Gary Geiss, Alexander Rosenberg, Charles Roco,

- Daniele Merico, Nima Alidoust, Hani Goodarzi, and Johnny Yu. Tahoe-100M : A Giga-Scale Single-Cell Perturbation Atlas for Context-Dependent Gene Function and Cellular Modeling. *bioRxiv (Cold Spring Harbor Laboratory) [preprint]*, 2025. doi: 10.1101/2025.02.20.639398. Preprint.
- [54] Dominik Klein, Giovanni Palla, Marius Lange, Michal Klein, Zoe Piran, Manuel Gander, Laetitia Meng-Papaxanthos, Michael Sterr, Lama Saber, Changying Jing, Aimée Bastidas-Ponce, Perla Cota, Marta Tarquis-Medina, Shrey Parikh, Ilan Gold, Heiko Lickert, Mostafa Bakhti, Mor Nitzan, Marco Cuturi, and Fabian J. Theis. Mapping cells through time and space with moscot. *Nature*, 638(8052):1065–1075, 2025. doi: 10.1038/s41586-024-08453-2.
- [55] Geoffrey Schiebinger, Jian Shu, Marcin Tabaka, Brian Cleary, Vidya Subramanian, Aryeh Solomon, Joshua Gould, Siyan Liu, Stacie Lin, Peter Berube, Lia Lee, Jenny Chen, Justin Brumbaugh, Philippe Rigollet, Konrad Hochedlinger, Rudolf Jaenisch, Aviv Regev, and Eric S. Lander. Optimal-transport analysis of single-cell gene expression identifies developmental trajectories in reprogramming. *Cell*, 176(4):928–943.e22, 2019. doi: 10.1016/j.cell.2019.01.006.
- [56] Marius Lange, Volker Bergen, Michal Klein, Manu Setty, Bernhard Reuter, Mostafa Bakhti, Heiko Lickert, Meshal Ansari, Janine Gote-Schniering, Herbert B. Schiller, Dana Pe’er, and Fabian J. Theis. CellRank for directed single-cell fate mapping. *Nature Methods*, 19(2):159–170, 2022. doi: 10.1038/s41592-021-01346-6.
- [57] Kexin Huang, Romain Lopez, Jan-Christian Hütter, Takamasa Kudo, Antonio Ríos, and Aviv Regev. Sequential Optimal Experimental Design of Perturbation Screens Guided by Multi-modal Priors. In *RECOMB 2024 (Lecture Notes in Computer Science)*, pages 17–37, 2024. doi: 10.1007/978-1-0716-3989-4_2.
- [58] Gerold Csendes, Gema Sanz, Kristóf Z. Szalay, and Bence Szalai. Benchmarking foundation cell models for post-perturbation RNA-seq prediction. *BMC Genomics*, 26(1):393–393, 2025. doi: 10.1186/s12864-025-11600-2.
- [59] Kasia Z. Kedzierska, Lorin Crawford, Ava P. Amini, and Alex X. Lu. Zero-shot evaluation reveals limitations of single-cell foundation models. *Genome biology*, 26(1):101–101, 2025. doi: 10.1186/s13059-025-03574-x.
- [60] Jialu Wu, Qing Ye, Yilin Wang, Renling Hu, Yiheng Zhu, Mingze Yin, Tianyue Wang, Jike Wang, Chang-Yu Hsieh, and Tingjun Hou. Biology-driven insights into the power of single-cell foundation models. *Genome biology*, 26(1):334–334, 2025. doi: 10.1186/s13059-025-03781-6.
- [61] Ramón Viñas, Maciej Wiatrak, Zoe Piran, Shuyang Fan, Liangze Jiang, Sarah A. Teichmann, Mor Nitzan, and Maria Brbić. Systema: a framework for evaluating genetic perturbation response prediction beyond systematic variation. *Nature Biotechnology*, 44(6):1050–1059, 2025. doi: 10.1038/s41587-025-02777-8.
- [62] Zhiting Wei, Yiheng Wang, Yicheng Gao, Shu-Guang Wang, Pu Li, Duanmiao Si, Yuli Gao, Siqi Wu, Danlu Li, Kejing Dong, Xingbo Yang, Chen Tang, Shaliu Fu, Xiaohan Chen, Wannian Li, Yuewen You, Chen Zhang, Aibin Liang, Guohui Chuai, and Qi Liu. Benchmarking algorithms for generalizable single-cell perturbation response prediction. *Nature Methods*, 23(2):451–464, 2025. doi: 10.1038/s41592-025-02980-0.

- [63] Malte D. Luecken, Scott Gigante, Daniel B. Burkhardt, Robrecht Cannoodt, Daniel Strobl, Nikolay S. Markov, Luke Zappia, Giovanni Palla, Wesley Lewis, Daniel Dimitrov, Michael Vinyard, D S Magruder, Michaela F. Mueller, Alma Andersson, Emma Dann, Qian Qin, Dominik J. Otto, Michal Klein, Olga Botvinnik, Louise Deconinck, Kai Waldrant, Sai Nirmayi Yasa, Artur Szalata, Andrew Benz, Zhijian Li, Open Problems Jamboree Members, Bastian Rieck, Constantin Ahlmann-Eltze, Eduardo da Veiga Beltrame, Carmen Bravo González-Blas, Ann T. Chen, Benjamin DeMeo, Can Ergen, Swann Floc’hlay, Adam Gayoso, Stephanie Hicks, Yuge Ji, Vitalii Kleshchevnikov, Gioele La Manno, Maximilian G. Lombardo, Romain Lopez, Dario Righelli, Hirak Sarkar, Valentine Svensson, Alexander Tong, Galen Xing, Chenling Xu, Jonathan M. Bloom, Angela Oliveira Pisco, Julio Saez-Rodriguez, Drausin Wulsin, Luca Pinello, Yvan Saeys, Fabian J. Theis, and Smita Krishnaswamy. Defining and benchmarking open problems in single-cell analysis. *Nature Biotechnology*, 43(7):1035–1040, 2025. doi: 10.1038/s41587-025-02694-w.
- [64] Xinjie Mao, Songming Zhang, Qianhong Wen, Xiangyu Wen, Kedu Jin, Hao Wu, Shuizhou Chen, Yuqiang Li, Lei Bai, Qi Liu, Ning Ding, Siqi Sun, and Zhangyang Gao. Benchmarking virtual cell models for in-the-wild perturbation response. *arXiv (Cornell University)*, 2026. doi: 10.48550/arxiv.2604.27646. Preprint.
- [65] Michael Vollenweider and Peter Bühlmann. Signal, Bounds, and Baselines: Principles for Evaluating Virtual Cell Perturbation Models. *bioRxiv (Cold Spring Harbor Laboratory)*, 2026. doi: 10.64898/2026.04.20.719650. Preprint. bioRxiv DOI on the openRxiv 10.64898 prefix (post-2025 migration); resolves via OpenAlex and the bioRxiv API. Title shown is the OpenAlex form; the bioRxiv record carries a revised subtitle across preprint versions (same DOI, authors, and posting date).
- [66] Ralf Schmidt, Zachary Steinhart, Madeline Layeghi, Jacob W. Freimer, Raymund Bueno, Vinh Nguyen, Franziska Blaesche, Chun Ye, and Alexander Marson. CRISPR activation and interference screens decode stimulation responses in primary human T cells. *Science*, 375(6580): eabj4008–eabj4008, 2022. doi: 10.1126/science.abj4008.

A Supplementary figures




Figure S1: Family-level method × capability matrix. Mean capability over each method family for six capabilities—measured-effect grounding, feasibility verdict, certificate, infeasibility proof, held-out validation, and combinatorial support—across the 92 surveyed methods (2011–2026). The reachability oracle (top row) is the only entry that combines full measured-effect grounding with a feasibility verdict, certificate, and infeasibility proof; every prior family that scores on the verdict/certificate columns scores near zero on measured-effect grounding, and vice versa.




Figure S2: Additivity is a bounded approximation across perturbation magnitudes. (a) For the 126 measured double perturbations in the Norman K562 screen, the cosine between the measured double and the sum of its two singles (median 0.71) does not collapse at large combined perturbation magnitude, but the scatter is real—epistasis exists. (b) The relative residual of the additive prediction is bounded and centered below one, so multi-gene recipes are an extrapolation to be tested at the bench rather than an exact prediction.




Figure S3: Positive-control ranking of T-helper regulators. Knockdown rank toward the T_H1 target (log scale; 1 = most aligned with the loss-of-function direction, 6871 = most anti-aligned) for canonical T_H1 and T_H2 regulators. T_H2 drivers, which must be removed, rank high (GATA3 at 155/6871); T_H1 drivers, which must be raised, rank low (TBX21 at 6775/6871). The transcription-factor-versus-marker split is an exploratory, post-hoc annotation, not a pre-registered test.




Figure S4: Triage funnel and the three actions a verdict enables. (a) The oracle compresses 102 ranked nominations to a testable set: 57 carry a druggable handle and 10 have a clinical-grade drug, while 45/102 (44%) are hard-to-drug. (b) Each verdict maps to one action—GO (reachable, run the minimal knockdown screen), REDIRECT (activation-required, switch to a CRISPRa arm), or STOP (provably outside, do not spend the arm). Every call is a falsifiable wet-lab hypothesis; the instrument triages, it does not validate.


DEG-weighted reachability verdict vs. unweighted

Figure S5: The verdict strengthens under DEG-magnitude weighting. For each target, the reachable cosine (**left**) and the held-out-gene z (**right**) under the default uniform metric (open markers) and under a DEG-magnitude weighting $w_g = |d_g|$ that concentrates the metric on the genes the target actually moves (filled markers). Reweighting *raises* the reachable cosine for every target—the Tier-2 $T_H2 \rightarrow T_H1$ shift rises from 0.627 to 0.803 at rest—and the held-out z stays far above the $z = 3$ significance floor, so the signal is not diluted by unchanged genes. This addresses the signal-dilution critique of Mejia et al. [14]; the unweighted metric remains the default and reproduces the published numbers exactly.




Figure S6: Each verdict is calibrated between its shuffled floor and an interpolated-duplicate ceiling. (left) For each target, the observed reachable cosine (unweighted, circle; DEG-weighted, diamond) placed between a shuffled-gene floor and an interpolated-duplicate ceiling (whiskers). Both floor and ceiling rise under weighting, so the observed value's placement between them is stable. (right) The scale-free dynamic-range fraction—the observed cosine expressed as a position between floor and ceiling—barely moves under weighting (Tier-2 rest 0.447 $\rightarrow$ 0.499), confirming that the strengthening in Fig. S5 is a genuine, calibrated signal rather than a scaling artifact. The dynamic-range calibration is computed on the rest condition (Monte-Carlo cost).




Figure S7: Cross-cell-type transfer: direction transfers, the minimal recipe does not. The unchanged operator applied to the Replogle et al. [15] K562 and RPE1 essential-gene screens (843 shared perturbations, 2,832 shared genes). (a) A single perturbation's effect *direction* transfers moderately (matched cross-type cosine median 0.35; grey band, shuffled-gene null $\pm 95\%$). (b) The reachability *verdict* transfers only as a graded residual: the median residual norm is systematically worse under a cross-type basis (hatched) than a within-type basis, for both K562 and RPE1 targets. (c) The minimal *recipe* does not transfer—same-target Jaccard overlap (0.11) collapses toward the random-selection null (0.05) under a cross-type basis. Direction is portable across cell types; the specific recipe is not.




Figure S8: The activation certificate is reproducible under a gene-axis split. (a) Mean rank of each top certificate gene when the certificate is refit on a random half of the gene axis (vertical bars, spread over 12 splits), against its published rank on the full cone fit (colour, fraction of half-fits keeping the gene in the top 30). The highest-demand genes (LYAR, CRTAM, GBP5, IRF8) hug the identity line; the tail (E2F3, LUZP1, SNTB2) drifts. (b) Rank agreement across independent halves (cross-half Spearman $\rho \approx 0.65$) and between a half-fit and the full fit ($\rho \approx 0.84$) is stable across 12 random split seeds and far above the shuffle null (grey band; cross-half $z \approx 35$).




Figure S9: Retrospective cross-reference of the top certificate genes tempers their biological reading. (a) The top 20 certificate genes ranked by unmet upward demand, coloured by literature direction-of-effect (SUPPORTS 4, PLAUSIBLE 7, UNKNOWN 9). The highest-demand genes are state-markers (LYAR, CRTAM, GBP5) and *negative* regulators (IKZF3/Aiolos, CBLB, LAG3), while the canonical Th1 master transcription factors (green stars) are present in the certificate but rank deep (IFNG #35, TBX21 #144, STAT4 #364, STAT1 #473). (b) Independent genetic corroboration: 8/20 genes carry a strong (≥ 0.5) Open Targets genetic-association score to a Th1-driven autoimmune disease. The certificate is best read as a reproducible ranking of unmet upward demand, not a validated list of activation targets; for the negative regulators, a useful intervention may be opposite in sign to naive activation. This cross-reference is exploratory and post-hoc.